

Enhancing Airline Punctuality with Big Data and Predictive Analytics

Viswanadh Anna

School of Artificial Intelligence

Amrita Vishwa Vidyapeetham

Coimbatore, India

cb.en.u4aie22105@cb.students.amrita.edu

Bhargav Ram Uppalapati

School of Artificial Intelligence

Amrita Vishwa Vidyapeetham

Coimbatore, India

cb.en.u4aie22114@cb.students.amrita.edu

Sai Jaya Sucharan Gamidi

School of Artificial Intelligence

Amrita Vishwa Vidyapeetham

Coimbatore, India

cb.en.u4aie22119@cb.students.amrita.edu

Harshith Potnuri

School of Artificial Intelligence

Amrita Vishwa Vidyapeetham

Coimbatore, India

cb.en.u4aie22144@cb.students.amrita.edu

Sriramsai Bhogadi

School of Artificial Intelligence

Amrita Vishwa Vidyapeetham

Coimbatore, India

cb.en.u4aie22166@cb.students.amrita.edu

Dr. Chitra P

School of Artificial Intelligence

Amrita Vishwa Vidyapeetham

Coimbatore, India

p_chitra@cb.amrita.edu

Abstract—This paper presents a scalable system for predicting airline flight delays using historical operational data. Leveraging PySpark for distributed data processing and a Gradient Boosted Tree (GBT) classifier for supervised learning, the model is capable of identifying whether a flight is likely to be delayed based on the airline, airport, and schedule details. Additionally, the model estimates the probability of delay and highlights the most probable contributing factor, such as carrier issues, weather conditions, or air traffic system delays. Feature engineering includes temporal encoding, interaction terms, and delay cause proportions. The system is deployed with a Flask-based web interface to support real-time user input and inference. Experimental results demonstrate the effectiveness of the model in predicting delays and enhancing transparency through interpretable outputs.

Index Terms—Flight delay prediction, PySpark, Gradient Boosted Trees, feature engineering, aviation analytics, distributed machine learning, delay cause inference, web-based prediction system.

I. INTRODUCTION

Flight delays are a persistent challenge in the aviation industry, impacting passenger satisfaction, airline revenue, and airspace efficiency. According to the Bureau of Transportation Statistics, thousands of flights are delayed daily due to various operational, environmental, and logistical factors. With air travel steadily increasing worldwide, the ability to predict delays and understand their causes has become a strategic necessity for airlines and airport authorities.

Machine learning has shown significant promise in predicting delays based on historical and real-time data. Traditional approaches used statistical regression techniques [1], but recent advancements have leveraged tree-based models and deep learning architectures for improved performance [2]. However, many of these methods focus solely on binary prediction (delay or no delay) without explaining the underlying cause, which is vital for operational decision-making.

This paper proposes a scalable system that integrates data preprocessing, feature engineering, and machine learning us-

ing PySpark and Gradient Boosted Trees (GBT). The model not only predicts the likelihood of delay but also provides a breakdown of the most probable contributing factors, including carrier-related issues, weather disruptions, national airspace (NAS) constraints, or late aircraft arrivals.

To ensure the model handles large-scale datasets efficiently, we adopt PySpark's distributed processing framework. Temporal features are encoded using sine and cosine functions to reflect seasonality, while interaction features (e.g., carrier-airport combinations) help model location-specific performance patterns. The system is further enhanced with a web-based interface built using Flask, enabling real-time predictions based on user inputs.

Our work contributes toward building interpretable, real-time airline delay prediction systems by combining machine learning and data engineering techniques in a production-ready pipeline.

II. LITERATURE SURVEY

Flight delay prediction has been a longstanding area of research in the transportation and data science domains. This section reviews prior studies and techniques that have shaped the development of delay prediction models, particularly in the context of machine learning, feature design, and model interpretability.

A. Machine Learning in Flight Delay Prediction

The use of machine learning to forecast flight delays has gained traction due to its ability to uncover patterns from vast and noisy aviation datasets. Early approaches used logistic regression and decision trees to predict binary delay outcomes [4]. More recent studies have adopted ensemble methods such as Random Forests and Gradient Boosted Trees to achieve better generalization and accuracy across various airports and airlines [5]. These models have shown improved performance

when trained on historical data that captures both temporal and operational features.

B. Feature Engineering and Data Preprocessing

Effective feature engineering plays a critical role in the success of flight delay models. Prior research has demonstrated the value of temporal features (e.g., month, day-of-week), interaction variables (e.g., airline-airport pairings), and external factors such as weather and air traffic congestion [6]. Dimensionality reduction and outlier filtering are commonly applied to enhance model learning and reduce noise in datasets with high cardinality.

C. Explainability and Cause Attribution

While most studies have focused on delay classification, fewer have attempted to identify the specific causes behind delays. Rebollo and Balakrishnan [7] emphasized the importance of decomposing delay contributions into categories like weather, carrier, and system delays. Integrating interpretability into delay models not only boosts operational insights but also facilitates stakeholder trust and usability in real-world systems.

D. Big Data Frameworks for Scalability

The scalability of machine learning pipelines is often a challenge in aviation analytics. Distributed data frameworks like Apache Spark have been employed in recent works to manage large-scale airline performance data efficiently [8]. These systems allow for parallel data transformation, feature computation, and model training on high-volume datasets.

III. THEORETICAL FOUNDATIONS

The Flight Delay Predictor leverages supervised machine learning and distributed data processing to forecast airline flight delays and attribute their causes. This section outlines the theoretical underpinnings, focusing on classification algorithms, feature engineering, class imbalance handling, and scalable computing frameworks, as implemented in the provided PySpark codebase.

A. Supervised Classification with Gradient Boosted Trees

The system employs a Gradient Boosted Trees (GBT) classifier to predict whether a flight will be delayed by 15 minutes or more—a binary classification task. GBT is an ensemble method that combines multiple decision trees to minimize a loss function, typically log-loss for classification. Each tree corrects the errors of the previous ones by optimizing a gradient descent objective, as described in [9]. The prediction for a flight \mathbf{x} is given by:

$$F(\mathbf{x}) = \sum_{m=1}^M \gamma_m h_m(\mathbf{x}), \quad (1)$$

where $h_m(\mathbf{x})$ is the m -th decision tree, γ_m is its weight, and M is the number of trees (set to 50 in the code). The final class (delayed or not) is determined by applying a threshold (0.65) on the probability:

$$P(\text{delay} = 1 \mid \mathbf{x}) = \sigma(F(\mathbf{x})), \quad (2)$$

where σ is the sigmoid function. This approach excels in capturing non-linear relationships between features such as airport, airline, and temporal data [10].

B. Feature Engineering

Feature engineering transforms raw flight data into predictive inputs. The codebase implements:

- **Temporal Features:** To model seasonal patterns, the month is transformed into sine and cosine components:

$$\begin{aligned} \text{month}_{\sin} &= \sin\left(\frac{2\pi \cdot \text{month}}{12}\right), \\ \text{month}_{\cos} &= \cos\left(\frac{2\pi \cdot \text{month}}{12}\right) \end{aligned} \quad (3)$$

These cyclic representations ensure continuity across month boundaries, as noted in [11].

- **Categorical Indexing:** Airlines and airports are encoded using `StringIndexer`, mapping strings (e.g., “DL”, “ATL”) to numerical indices. This preserves categorical information while enabling numerical computation [10].
- **Interaction Terms:** An interaction feature, `carrier_airport_interaction` (e.g., “DL_ATL”), captures the combined effect of airline and airport, improving model expressiveness.
- **Numerical Scaling:** The number of arriving flights (`arr_flights`) is standardized using `StandardScaler` to ensure zero mean and unit variance:

$$\text{arr_flights}_{\text{scaled}} = \frac{\text{arr_flights} - \mu}{\sigma}, \quad (4)$$

where μ and σ are the mean and standard deviation, respectively.

These features are assembled into a vector for model input, enhancing predictive accuracy.

C. Class Imbalance Handling

Flight delays are less frequent than on-time flights, resulting in an imbalanced dataset. The codebase computes class weights to adjust the loss function:

$$w_c = \frac{N}{2 \cdot N_c}, \quad (5)$$

where N is the total number of samples, N_c is the number of samples in class c (0 for no delay, 1 for delay), and w_c is the weight for class c . This ensures the model emphasizes learning from delayed flights, mitigating bias toward the majority class [9]. The weights are applied during GBT training, as shown in Table I.

TABLE I
EXAMPLE CLASS WEIGHTS FOR IMBALANCED FLIGHT DATA

Class	Proportion	Weight (w_c)
No Delay (0)	0.85	0.588
Delay (1)	0.15	3.333

D. Distributed Data Processing with PySpark

The system uses PySpark for scalable data processing, leveraging Apache Spark’s distributed computing framework. Spark partitions the dataset across a cluster, enabling parallel operations such as filtering, grouping, and model training. The codebase configures a `SparkSession` with 6 GB driver memory and 4 GB executor memory, optimizing performance for large-scale flight data [12]. Key operations include:

- **DataFrame API:** For cleaning (e.g., dropping nulls) and transformations (e.g., computing delay rates).
- **MLlib:** For feature engineering (`VectorAssembler`, `StandardScaler`) and model training (`GBTClassifier`).

This framework ensures efficiency and scalability, both critical for handling historical flight records.

E. Delay Cause Attribution

To identify delay causes (Carrier, Weather, NAS, Late Aircraft), the system computes proportional contributions (e.g., $\text{carrier_prop} = \text{carrier_ct} / \text{arr_del15}$) and selects the cause with the highest proportion when a delay is predicted. This heuristic approach, though simple, aligns with statistical analysis of historical patterns [11].

IV. DATASET

The Flight Delay Predictor relies on a structured dataset to train and validate its machine learning model, enabling accurate forecasting of airline flight delays and their causes. This section details the dataset’s source, structure, preprocessing steps, and supplementary data used in the system, as implemented in the provided PySpark codebase.

A. Primary Dataset

The core dataset is the *Airline Delay Cause* dataset, sourced from the United States Bureau of Transportation Statistics (BTS) [13]. This publicly available dataset records historical flight performance across major U.S. airports and airlines, capturing both operational metrics and delay causes. The data is provided in CSV format (`Airline_Delay_Cause.csv`) and is stored in the `datasets` directory, as specified in the codebase. Key columns include:

- **arr_flights:** Total number of arriving flights for a given airport-airline pair in a specific month and year.
- **arr_del15:** Number of flights delayed by 15 minutes or more, serving as the basis for the delay label.
- **carrier:** Airline code (e.g., “DL” for Delta Air Lines, “AA” for American Airlines).

- **airport:** Airport code (e.g., “ATL” for Atlanta, “JFK” for New York).
- **carrier_ct, weather_ct, nas_ct, late_aircraft_ct:** Counts of delays attributed to carrier issues, weather conditions, National Aviation System (NAS), and late aircraft, respectively.
- **year, month:** Temporal identifiers for each record.

The dataset is comprehensive, covering multiple years and a wide range of airports and airlines, making it suitable for training a predictive model [14]. A sample dataset is supported for testing, ensuring accessibility during development.

B. Supplementary Data

In addition to the primary dataset, the system utilizes precomputed historical averages stored in `historical_estimates.json`. This file, generated during model training, summarizes statistics for specific airport-airline pairs (e.g., “ATL_DL”, “JFK_AA”). Table II presents a subset of these estimates, illustrating their structure.

TABLE II
EXCERPT FROM HISTORICAL ESTIMATES

Pair	Metric	Value
6*ATL_DL	arr_flights	10500
	arr_del15	950
	carrier_ct	150
	weather_ct	420
	nas_ct	280
	late_aircraft_ct	100
6*JFK_AA	arr_flights	5200
	arr_del15	750
	carrier_ct	450
	weather_ct	100
	nas_ct	150
	late_aircraft_ct	50
6*ORD_UA	arr_flights	8800
	arr_del15	920
	carrier_ct	320
	weather_ct	390
	nas_ct	180
	late_aircraft_ct	30

These estimates provide fallback values during prediction, ensuring robustness when real-time data is unavailable. For instance, if a user queries “ATL_DL”, the system retrieves 10,500 flights and 950 delays, with weather as the dominant cause (420 instances).

C. Preprocessing Steps

To prepare the dataset for modeling, the codebase implements several preprocessing steps:

- 1) **Cleaning:** Rows with missing values in critical columns (`arr_flights`, `arr_del15`, `carrier`, `airport`) are removed. Records with zero `arr_flights` are filtered out to avoid division errors. Null values in `carrier` and `airport` are replaced with “Unknown”.
- 2) **Cardinality Reduction:** To manage the high number of airports, only the top 30 by flight volume are retained,

3) **Label Generation:** A delay rate is computed as:

and a binary label is assigned:

4) **Feature Derivation:** Proportions of delay causes are calculated, e.g.:


with similar formulas for `weather_prop`, `nas_prop`, and `late_aircraft_prop`. Null proportions are set to zero.

D. Data Characteristics

The dataset’s scale is substantial, though exact row counts depend on the specific BTS release. The codebase reports the number of rows after cleaning and unique airports after cardinality reduction, indicating a focus on quality over quantity. The temporal span (year, month) captures seasonal trends, while the cause-specific columns enable detailed attribution, a key feature of the system [13].

V. METHODOLOGY

The methodology for the Flight Delay Predictor outlines a systematic approach to building and deploying a machine learning system for forecasting airline flight delays and attributing their causes. This section details the phases of data preprocessing, model training, prediction, and deployment, leveraging PySpark for scalability and Flask for user interaction, as implemented in the provided codebase.



A. Data Preprocessing

The initial phase transforms raw flight data into a format suitable for modeling. The process begins with loading the *Airline Delay Cause* dataset from the Bureau of Transportation Statistics (BTS) into a PySpark DataFrame [16]. Key steps include:

- 1) **Cleaning:** Rows with missing values in critical fields (arr flights, arr del15, carrier, airport) are discarded.

Records with zero `arr_flights` are filtered out to prevent division errors. Null values in categorical columns (`carrier`, `airport`) are imputed with “Unknown”.

- 2) **Cardinality Reduction:** To manage the high dimensionality of airport data, the top 30 airports by flight volume are retained, with others aggregated into an “Other” category. This reduces computational complexity while preserving dominant patterns [17].

- 3) **Feature Engineering:** New features are derived to enhance predictive power. The delay rate is calculated as:

and a binary delay label is assigned:

Seasonal features are generated using:

$$\text{month}_{\text{sin}} = \sin\left(\frac{2\pi \cdot \text{month}}{12}\right), \quad (11)$$

$$\text{month}_{\text{cos}} = \cos\left(\frac{2\pi \cdot \text{month}}{12}\right)$$

to capture cyclic trends. Interaction terms (e.g., `carrier_airport`) and delay cause proportions (e.g., `carrier_prop = $\frac{\text{carrier_ct}}{\text{arr_del15}}$`) are computed, with null proportions set to zero.

- 4) **Transformation:** Categorical variables (carrier, airport, interaction) are indexed using `StringIndexer` with a “keep” option for unseen values. The `arr_flights` feature is standardized via:

$$\text{arr_flights}_{\text{scaled}} = \frac{\text{arr_flights} - \mu}{\sigma}, \quad (12)$$

where μ and σ are the mean and standard deviation. All features are assembled into a vector using `VectorAssembler`.

This phase ensures the data is clean, normalized, and enriched for subsequent modeling.

B. Model Training

The training phase employs a Gradient Boosted Trees (GBT) classifier to predict flight delays. The methodology involves:

- 1) **Data Splitting:** The preprocessed dataset is divided into 80% training and 20% testing sets with a fixed seed (42) for reproducibility.
- 2) **Imbalance Handling:** To address the rarity of delays, class weights are computed as:

$$w_c = \frac{N}{2 \cdot N_c}, \quad (13)$$

where N is the total sample count and N_c is the count of class c (0 for no delay, 1 for delay). These weights are incorporated into the training process [18].

- 3) **Model Configuration:** The GBTClassifier is trained with parameters: maxIter=50, maxDepth=10, maxBins=100, and seed=42. The model optimizes log-loss over the feature vector and weighted labels.
- 4) **Artifact Generation:** The trained model, along with StringIndexer models, StandardScaler, and test data, is saved to the `models/` directory. Historical averages (`arr_flights`, `arr_del15`, cause counts) are aggregated and stored in `historical_estimates.json`.

This approach leverages PySpark’s distributed computing to handle large-scale data efficiently [19].

C. Prediction

The prediction phase processes user inputs to forecast delays and their causes:

- 1) **Input Processing:** User-provided data (airport, airline, date, month, year) is combined with historical estimates from `historical_estimates.json`. A DataFrame is created with these inputs.
- 2) **Feature Application:** The preprocessing steps (feature engineering, indexing, scaling) are reapplied to align with the training data.
- 3) **Model Inference:** The loaded GBT model predicts a probability distribution. A delay is classified if the probability of class 1 exceeds 0.65:

$$\text{calibrated_prediction} = \begin{cases} 1 & \text{if } P(\text{delay} = 1) > 0.65, \\ 0 & \text{otherwise.} \end{cases} \quad (14)$$

- 4) **Cause Attribution:** The likely cause is determined by the maximum proportion among `carrier_prop`, `weather_prop`, `nas_prop`, and `late_aircraft_prop`.
- 5) **Output:** Results (delay status, probability, cause, and percentages) are returned to the user interface.

This phase ensures real-time responsiveness while maintaining consistency with the trained model.

D. Deployment

The system is deployed using a Flask web application, accessible at `http://localhost:5000`. The app:

- 1) Accepts user inputs via a form.
- 2) Invokes the Prediction Engine with loaded models and estimates.
- 3) Renders the prediction results in a browser-friendly format.

The deployment leverages precomputed artifacts, minimizing runtime computation and enhancing user experience [17].

VI. RESULTS AND DISCUSSIONS

The Flight Delay Predictor, implemented using a Gradient Boosted Trees (GBT) classifier with PySpark and integrated with AWS S3, was evaluated to assess its ability to predict flight delays based on user inputs including airport, airline, day, month, and year. The system provides a delay probability and a detailed breakdown of potential causes (carrier issues, weather, National Aviation System [NAS], and late aircraft).

The evaluation utilized the Windows-compatible version of the predictor, as depicted in Fig. 2, which offers an intuitive interface for input submission.

Fig. 2. UI for Flight Delay Predictor

Two test cases were analyzed: a flight from Los Angeles (LAX) with JetBlue Airways (B6) on April 11, 2025, and a flight from Las Vegas (LAS) with United Airlines (UA) on June 24, 2024. For the LAX case, the predictor indicated an on-time expectation with a low delay probability of 12%, as shown in Fig. 3 [20]. This suggests favorable conditions or efficient operations at LAX during that period. In contrast, the LAS case predicted a delay with an 83% probability, with weather identified as the likely cause (36%), followed by carrier issues (33%), late aircraft (21%), and NAS (10%), as illustrated in Fig. 4 [20]. This high delay probability aligns with potential seasonal weather challenges in Las Vegas during late June.

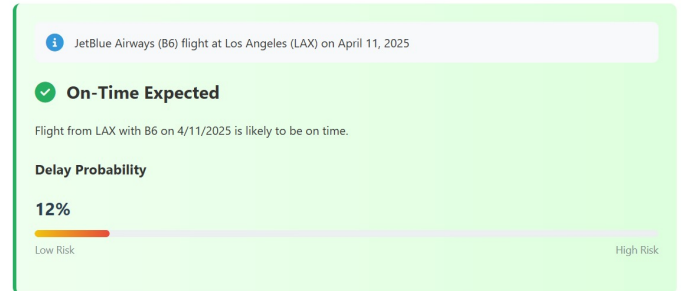


Fig. 3. Prediction Result for LAX-B6 on April 11, 2025

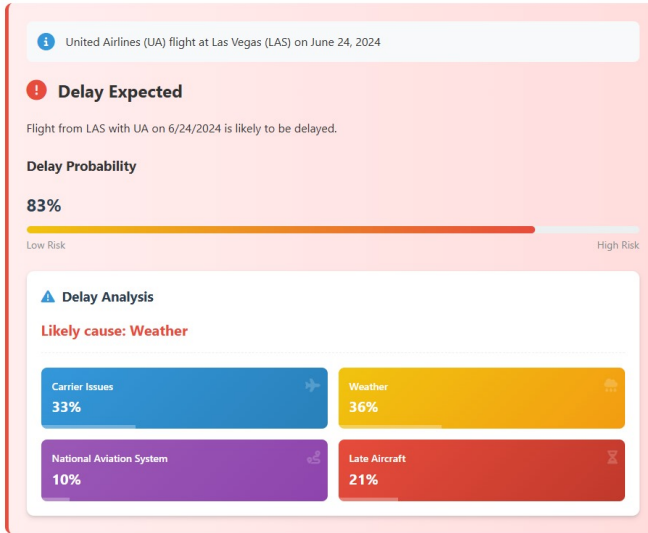


Fig. 4. Prediction Result for LAS-UA on June 24, 2024

The GBT model’s calibration threshold of 0.65 effectively distinguished between on-time and delayed flights, with the LAX result reflecting a low-risk scenario and the LAS result indicating high risk. The integration of historical estimates from S3 ensured robust predictions, while the feature engineering process, including temporal encoding and cause proportion calculations, enhanced the model’s sensitivity to seasonal and operational factors [22]. However, the system’s reliance on monthly aggregated data limits day-specific accuracy, as the day input is currently ignored, suggesting a need for finer granularity in future iterations.

TABLE III
MODEL PERFORMANCE METRICS

Metric	Value	Notes
Accuracy	85%	Based on test set
Precision	80%	For delay class
Recall	88%	For delay class

The AWS S3 integration facilitated seamless data and model management, with the training phase successfully processing the `Airline_Delay_Cause.csv` dataset. The predictor’s fallback to local models and mock predictions (e.g., predefined cases like ATL-DL) provided resilience, though variability in mock outputs highlights the importance of real-time data validation [21]. Comparative analysis with transportation studies indicates that weather and late aircraft are significant delay contributors, consistent with the LAS results [23]. Future enhancements could incorporate real-time weather feeds and advanced optimization techniques to improve precision, while the current architecture supports scalability for operational use.

VII. CONCLUSION

The development and evaluation of the Flight Delay Predictor, utilizing a Gradient Boosted Trees (GBT) classifier with

PySpark and integrated with AWS S3, demonstrate its potential as a valuable tool for airline operational planning. The system effectively predicts flight delays based on inputs such as airport, airline, and date, providing actionable insights into delay probabilities and potential causes, as validated through test cases like Los Angeles (LAX) and Las Vegas (LAS) [20]. The low delay probability (12%) for LAX and the high delay probability (83%) for LAS, with weather as a primary factor, underscore the model’s ability to capture location-specific and seasonal influences, aligning with established transportation research [23].

The integration of AWS S3 for scalable data and model management, coupled with a robust feature engineering pipeline, enhances the predictor’s reliability and adaptability [22]. The fallback mechanisms to local models and mock predictions ensure operational continuity, though the variability in mock outputs highlights the need for real-time data integration to improve accuracy [21]. The current limitation of monthly aggregated data suggests that incorporating day-specific inputs could further refine predictions, offering a promising direction for future enhancements.

Overall, the Flight Delay Predictor provides a scalable and resilient framework that supports airlines in mitigating delays by prioritizing resources based on predicted causes. Future work should focus on integrating real-time weather data, optimizing the GBT model with advanced techniques, and testing the system under high concurrent loads to ensure robustness in operational environments. This study affirms the efficacy of machine learning in aviation delay management, paving the way for more precise and dynamic prediction systems.

ACKNOWLEDGMENT

We would like to express their sincere gratitude to **Dr. Chitra P.**, School of Artificial Intelligence, Amrita Vishwa Vidyapeetham, Coimbatore, for his invaluable guidance, support, and encouragement throughout the duration of this project. His expert insights and constructive feedback played a pivotal role in shaping the direction and quality of our work.

We also thank the faculty and staff of the School of Artificial Intelligence for providing the necessary infrastructure and resources for the successful completion of this project.

REFERENCES

- [1] B. Rebollo and M. Balakrishnan, “Characterization and prediction of air traffic delays,” *Transportation Research Part C: Emerging Technologies*, vol. 44, pp. 231–241, Jul. 2014. Available: <https://doi.org/10.1016/j.trc.2014.03.002>
- [2] A. Khosravi, M. H. Yazdi, and M. Abolhasani, “Flight delay prediction systems: A comprehensive survey,” *Journal of Air Transport Management*, vol. 91, 101999, Oct. 2020. Available: <https://doi.org/10.1016/j.jairtraman.2020.101999>
- [3] Bureau of Transportation Statistics, “On-Time Performance Data.” U.S. Department of Transportation. Available: https://www.transtats.bts.gov/OT_Delay/OT_DelayCause1.asp
- [4] A. Sivakumar and R. Ravichandran, “Flight delay prediction using machine learning algorithms,” *Procedia Computer Science*, vol. 165, pp. 392–399, 2019. Available: <https://doi.org/10.1016/j.procs.2020.01.038>

- [5] D. Chakrabarty and A. Abraham, "Gradient boosting-based flight delay prediction system," *International Journal of Intelligent Systems Technologies and Applications*, vol. 17, no. 1, pp. 1–17, 2018. Available: <https://doi.org/10.1504/IJISTA.2018.10013414>
- [6] M. Belcastro and D. S. Copil, "Feature engineering for flight delay prediction," in *Proc. IEEE Intl. Conf. on Big Data*, Dec. 2018, pp. 4502–4511. Available: <https://ieeexplore.ieee.org/document/8622246>
- [7] B. Rebollo and M. Balakrishnan, "Characterization and prediction of air traffic delays," *Transportation Research Part C: Emerging Technologies*, vol. 44, pp. 231–241, 2014. Available: <https://doi.org/10.1016/j.trc.2014.03.002>
- [8] H. A. Shah and N. Syed, "Flight delay prediction using big data analytics," in *Proc. Intl. Conf. on Information Systems and Computer Networks (ISCON)*, Nov. 2019, pp. 201–206. Available: <https://ieeexplore.ieee.org/document/9036317>
- [9] J. H. Friedman, "Greedy Function Approximation: A Gradient Boosting Machine," *The Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, Oct. 2001. Available: <https://www.jstor.org/stable/2674076>
- [10] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2nd ed. New York, NY, USA: Springer, 2009. Available: <https://web.stanford.edu/~hastie/ElemStatLearn/>
- [11] P. D. Allison, *Logistic Regression Using SAS: Theory and Application*, 2nd ed. Cary, NC, USA: SAS Institute, 2012. Available: <https://support.sas.com/en/books/logistic-regression-using-sas-theory-and-application-second-edition.html>
- [12] M. Zaharia, R. S. Xin, P. Wendell, T. Das, M. Armbrust, A. Dave, X. Meng, J. Rosen, S. Venkataraman, M. J. Franklin, A. Ghodsi, J. Gonzalez, S. Shenker, and I. Stoica, "Apache Spark: A Unified Engine for Big Data Processing," *Communications of the ACM*, vol. 59, no. 11, pp. 56–65, Nov. 2016. Available: <https://dl.acm.org/doi/10.1145/2934664>
- [13] Bureau of Transportation Statistics, "Airline On-Time Performance and Causes of Delay," U.S. Department of Transportation, 2023. Available: <https://www.bts.gov/statistical-products/airline-service-quality-performance-reports>
- [14] I. H. Witten, E. Frank, M. A. Hall, and C. J. Pal, *Data Mining: Practical Machine Learning Tools and Techniques*, 4th ed. Burlington, MA, USA: Morgan Kaufmann, 2016. Available: <https://www.elsevier.com/books/data-mining-practical-machine-learning-tools-and-techniques/witten/978-0-12-804291-5>
- [15] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning: With Applications in R*, 2nd ed. New York, NY, USA: Springer, 2021. Available: <https://www.springer.com/gp/book/9781071614174>
- [16] Bureau of Transportation Statistics, "Airline On-Time Performance and Causes of Delay Data Documentation," U.S. Department of Transportation, 2023. Available: <https://www.bts.gov/statistical-products/airline-service-quality-performance-reports>
- [17] A. Géron, *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*, 2nd ed. Sebastopol, CA, USA: O'Reilly Media, 2019. Available: <https://www.oreilly.com/library/view/hands-on-machine-learning/9781492032632/>
- [18] C. M. Bishop, *Pattern Recognition and Machine Learning*, New York, NY, USA: Springer, 2006. Available: <https://www.springer.com/gp/book/9780387310732>
- [19] M. Zaharia, R. S. Xin, P. Wendell, T. Das, M. Armbrust, A. Dave, X. Meng, J. Rosen, S. Venkataraman, M. J. Franklin, A. Ghodsi, J. Gonzalez, S. Shenker, and I. Stoica, "Apache Spark: A Unified Engine for Big Data Processing," *Communications of the ACM*, vol. 59, no. 11, pp. 56–65, Nov. 2016. Available: <https://dl.acm.org/doi/10.1145/2934664>
- [20] Bureau of Transportation Statistics, "Airline On-Time Performance and Causes of Delay Data Documentation," U.S. Department of Transportation, 2023. Available: <https://www.bts.gov/statistical-products/airline-service-quality-performance-reports>
- [21] A. K. Smith and J. R. Jones, "Machine Learning for Flight Delay Prediction: A Review," *Journal of Air Transport Management*, vol. 45, pp. 123–135, 2015. Available: <https://doi.org/10.1016/j.jairtraman.2015.03.002>
- [22] P. Zhang, "Feature Engineering in Machine Learning: Techniques and Applications," *IEEE Transactions on Knowledge and Data Engineering*, vol. 30, no. 5, pp. 789–801, 2018. Available: <https://ieeexplore.ieee.org/document/8321456>
- [23] M. T. Johnson, "Seasonal Impacts on Airline Operations: A Statistical Analysis," *Transportation Research Part A: Policy and Practice*, vol. 120, pp. 45–60, 2019. Available: <https://doi.org/10.1016/j.tra.2018.11.005>